

TERRYTORY

Strategic Review

Where the product stands, what it is worth, and how it could plausibly pay for itself on a EUR 50 monthly ceiling.

PREPARED FOR

Lukáš — Founder & CEO

SCOPE

Product · Positioning · Monetisation

STATUS

Pre-revenue · Pre-audience

You have built a factory, not a magazine

The engineering is genuinely strong. The publication it produces barely exists yet. That gap is the whole story, and it is a far better problem to have than the reverse.

2

articles live

165

stories queued

0

analytics installed

~20

EUR/mo run cost

Territory ingests roughly 130 tech and AI stories a day, scores each one, proposes article angles, writes them in a chosen format, and publishes on a schedule without supervision. That pipeline works end to end. It ran unattended last week and produced 93 usable story candidates in a single pass.

Against that: **two articles are actually live**. The tech publication has published nothing at all. There is no analytics on the site, so traffic is not merely low, it is *unmeasured*. The bottleneck is not the machine. It is that the machine's output has not been shipped.

The one-line read. Six months of engineering has produced a capable content system and almost no content. The next 90 days should be spent shipping and measuring, not building. Revenue in that window will be near zero, and that is the correct expectation — the asset being built is an audience, and audiences are slow.

What actually exists

Six stages, fully automated, from raw feed to live page.

STAGE	WHAT IT DOES
1. Ingest	Five sources per vertical (MIT Tech Review, VentureBeat, Import AI, The Verge, Ars Technica, TechCrunch, Wired, Engadget). Deduplicated on a rolling hash so nothing is processed twice.
2. Score	Every story is rated 1–10 by a cheap model against a consumer-appeal rubric. Only 8.0+ enters the pool.
3. Plan	A 'Creative Director' agent reads the day's top stories and proposes 2–5 article concepts — each with a format, an angle, a target reader, and a stated reason it matters today.
4. Write	Chosen concepts get web research for factual grounding, then drafting in one of six formats (article, hot take, explainer, listicle, roundup, comparison) under a named author persona.
5. Edit	A full review surface: inline rich-text editing, headline variants, image slots, social embeds, SEO fields.
6. Ship	Scheduled publishing runs in the cloud every fifteen minutes, so articles go live whether or not any machine of yours is on.

Built in the last cycle

- **Editorial memory.** The system now records every change you make to a draft — headline rewrites, cut paragraphs, deleted punctuation tics — and feeds those patterns back into the next article's prompt. Your taste compounds instead of evaporating.
- **Vacation mode.** A single switch that halts every paid process. Currently engaged.
- **Security.** The entire admin API was publicly readable, including unpublished drafts. Closed and verified in production.
- **Test coverage.** From zero to seventeen smoke tests running on every push.

What is genuinely good here

The editorial memory loop

This is the most commercially interesting thing you own. Most AI content operations produce generic output forever. Yours gets closer to your voice with every article you touch. That is a real moat, and it is also the single best defence against being classified as low-value AI content.

Format intelligence

The planner does not simply write 'an article'. It decides that a given story wants to be a hot take, and a different one wants to be a comparison. That is an editorial judgement most competitors do not attempt.

Genuine autonomy

Scraping, scoring, planning and publishing all run in the cloud on schedules. The system does not need you awake. Very few solo operations reach this point.

Multi-vertical from day one

Adding a third vertical is a config file, not a rewrite. If one niche finds traction, replication is nearly free.

Cost control

Model tiers, a spend ceiling, and a hard off switch are already built. You can throttle to near zero without touching code.

Worth stating plainly. The codebase is more valuable than the website. If Terrytory never earns a euro from advertising, the system itself — a working, cost-controlled, self-improving publishing engine — is a portfolio piece and a saleable product in its own right.

What is holding it back

Ordered by how much each one costs you. The first three are severe.

CRITICAL You are flying blind

There is no analytics on the site. No Google Analytics, no Plausible, nothing. Your stated goal is to see traffic — and right now you structurally cannot, even if it existed. This is a thirty-minute fix and it blocks every other decision.

CRITICAL The publication is nearly empty

Two live articles on AI. Zero on tech. A magazine with two articles will not rank, will not be approved for advertising, and gives a first-time visitor no reason to return. Volume is the prerequisite for everything else.

CRITICAL The funnel is jammed at your desk

165 scored stories and 36 planned concepts are sitting unused. The machine's throughput vastly exceeds the rate at which drafts get approved. Everything upstream is working and nothing is coming out the end.

HIGH No images on any published article

Zero of fourteen published pieces carry a hero image. For a publication positioned as a magazine, this is the most visible quality gap on the site.

HIGH Subdomains do not resolve

ai.territory.xyz and tech.territory.xyz have no DNS. Everything is served from a single www address, so the two publications have no separate identity despite the code fully supporting it.

MEDIUM The scoring rubric may be too generous

In the last run, 72 of 95 tech stories cleared a threshold nominally set to admit the top fifth. Either the vertical had an extraordinary day or the grader is lenient. If lenient, you are paying to write mediocre stories.

MEDIUM The admin cannot save on the live site

The current host has a read-only filesystem, so most editing actions silently fail in production and work only on your own machine. A migration to a normal server removes this entire class of problem.

Where this is heading

Read the commit history and the direction is unmistakable. The system started as a scraper that produced articles. It is becoming an autonomous newsroom with you as editor-in-chief rather than writer.

PHASE	STATUS	DETAIL
Phase 1 Scrape and publish	Done	Raw pipeline. Feed in, article out.
Phase 2 Editorial judgement	Done	Scoring, format selection, the Creative Director proposing angles rather than just summarising.
Phase 3 Taste transfer	Just built	The system learns your edits. This is the shift from 'AI writes, you fix' toward 'AI writes the way you would have'.
Phase 4 Distribution	Not started	SEO, newsletter, social. Currently the largest gap and the one that decides whether any of this reaches a reader.
Phase 5 Monetisation	Not started	Only rational once Phase 4 produces measurable traffic.

The strategic fork

There are two honest futures, and they need different behaviour from you.

A publication

You are building an audience. Success is measured in returning readers and search rankings. Requires consistency for months before anything happens. Monetises via affiliate links, then advertising, eventually sponsorship.

A product

The engine is the asset. Other people want an automated niche publication and cannot build one. Monetises far faster than advertising, at meaningfully higher value per customer. Territory becomes the demo.

You do not have to choose today. Both paths require the same next ninety days: publish consistently and measure what happens. But it is worth knowing that the second path is where the money realistically is, and the first path is what makes the second one credible.

How this could actually make money

You asked directly about advertising and subscriptions. The honest answer is that neither works yet, for reasons worth understanding rather than guessing at.

CHANNEL	VERDICT	REASONING
Display advertising	Not yet	Google AdSense requires approval, and AI-assisted content sites are refused more often than not — particularly thin ones. Even once approved, tech content earns roughly EUR 3–8 per thousand page views. Your EUR 3–4 target therefore needs about a thousand monthly readers. That is achievable, but it is a Q4 conversation, not a July one.
Affiliate links	Start here	The single best fit for your situation, and the infrastructure is already in the codebase — the content config supports affiliate slots that the writer can place mid-article. Crucially, affiliate works at low traffic: one reader buying a EUR 200 device can out-earn ten thousand ad impressions. Tech is close to the ideal vertical for this.
Subscriptions	No	Readers subscribe to people and to information they cannot get elsewhere. Aggregated tech news is neither. Pursuing this now would consume months and convert nobody.
Newsletter	Start capturing	Not directly monetisable at small scale, but it is the only asset here that you own outright rather than rent from Google. Costs nothing to begin, compounds from day one, and becomes sponsorable at around a thousand subscribers. Add the signup form before the traffic arrives, not after.
Selling the engine	Highest value	Discussed opposite. Materially faster to revenue than advertising, and Territory itself is the proof it works.

The risk nobody mentions. Google's spam policy explicitly targets content produced at scale primarily to rank. A fully autonomous site publishing multiple AI-written articles daily fits that description precisely, and sites matching it have been removed from search results entirely. Your mitigation is the editorial memory system: genuinely edited, opinionated content with a consistent human voice is a different category. This is a strong argument for **one carefully edited article a day rather than eight automatic ones**.

Living inside EUR 50 a month

This is comfortably achievable. Your recent spend of roughly EUR 3 a day came almost entirely from grading every one of 130 stories daily with a paid model. That is the lever.

LINE ITEM	EUR/MO	NOTE
Server Hetzner CX22 or similar	4–6	Cheaper and better suited than most managed hosts. Removes the read-only filesystem problem entirely.
Domain amortised annually	~1	Already owned.
Story grading reduced to top 30/day	3–6	The main saving. Grade the best thirty headlines, not all 130.
Research 2 articles/day	1–2	Perplexity Sonar, cheap.
Writing 2 articles/day, low tier	2–4	DeepSeek is close to free at this volume.
Images 2–3 per article	4–8	The largest variable cost. Worth capping.
Total	15–27	Roughly half your ceiling, leaving room to experiment.

Three rules to stay inside it

- **Grade fewer stories.** The headline pre-ranking already exists and is nearly free. Only send the top thirty for full paid scoring.
- **One article per vertical per day.** Not eight. Better for cost, better for quality, and materially safer with search engines.
- **Cap image generation.** Two per article, and reuse where sensible. This is the least predictable line and the easiest to let run away.

The next ninety days

Build nothing new. Everything below is either configuration, publishing, or measurement.

This week

- **Install analytics.** Plausible or Google Analytics, plus Google Search Console. Half an hour. Without this, nothing else can be judged.
- **Point the subdomains at the site.** The code already routes them; only DNS is missing.
- **Publish five tech articles.** The vertical currently has none. Eighty-one scored stories are waiting.
- **Turn the spend tracker on** before resuming automation, so the first bill is a confirmation rather than a surprise.

This month

- **Move off the read-only host.** It eliminates an entire class of bug and cuts the monthly cost.
- **One article per vertical per day, every day.** Consistency matters more to search engines than volume.
- **Edit every single one by hand.** Not for quality alone — every edit now trains the system. Early edits are worth more than later ones.
- **Fix the images.** Every published article should carry one.
- **Add the newsletter signup,** even with no readers yet.

This quarter

- **Reach sixty published articles.** Roughly the threshold where search engines begin treating a site as a real publication.
- **Enable affiliate slots** on articles where a recommendation is genuinely useful. Never where it is not.
- **Apply for AdSense** only once traffic is real and the archive is substantial.
- **Review whether the engine is the product.** By then you will have the data to judge it.

The realistic ninety-day outcome. Around sixty articles published, somewhere between one and five hundred monthly readers if the SEO lands, and plausibly zero to five euro of revenue. That is not a disappointing result — it is what the first quarter of a publication looks like. What you will actually own is a working system, real traffic data, and a trained editorial model that no competitor can copy. The money question becomes answerable in October, not July.

Prepared for the founder. Figures drawn from the live repository and the most recent unattended production run.